

ORACLE

A 3D network diagram showing a central Earth globe connected by thin lines to several surrounding silver spherical nodes, representing a global network or cloud architecture.

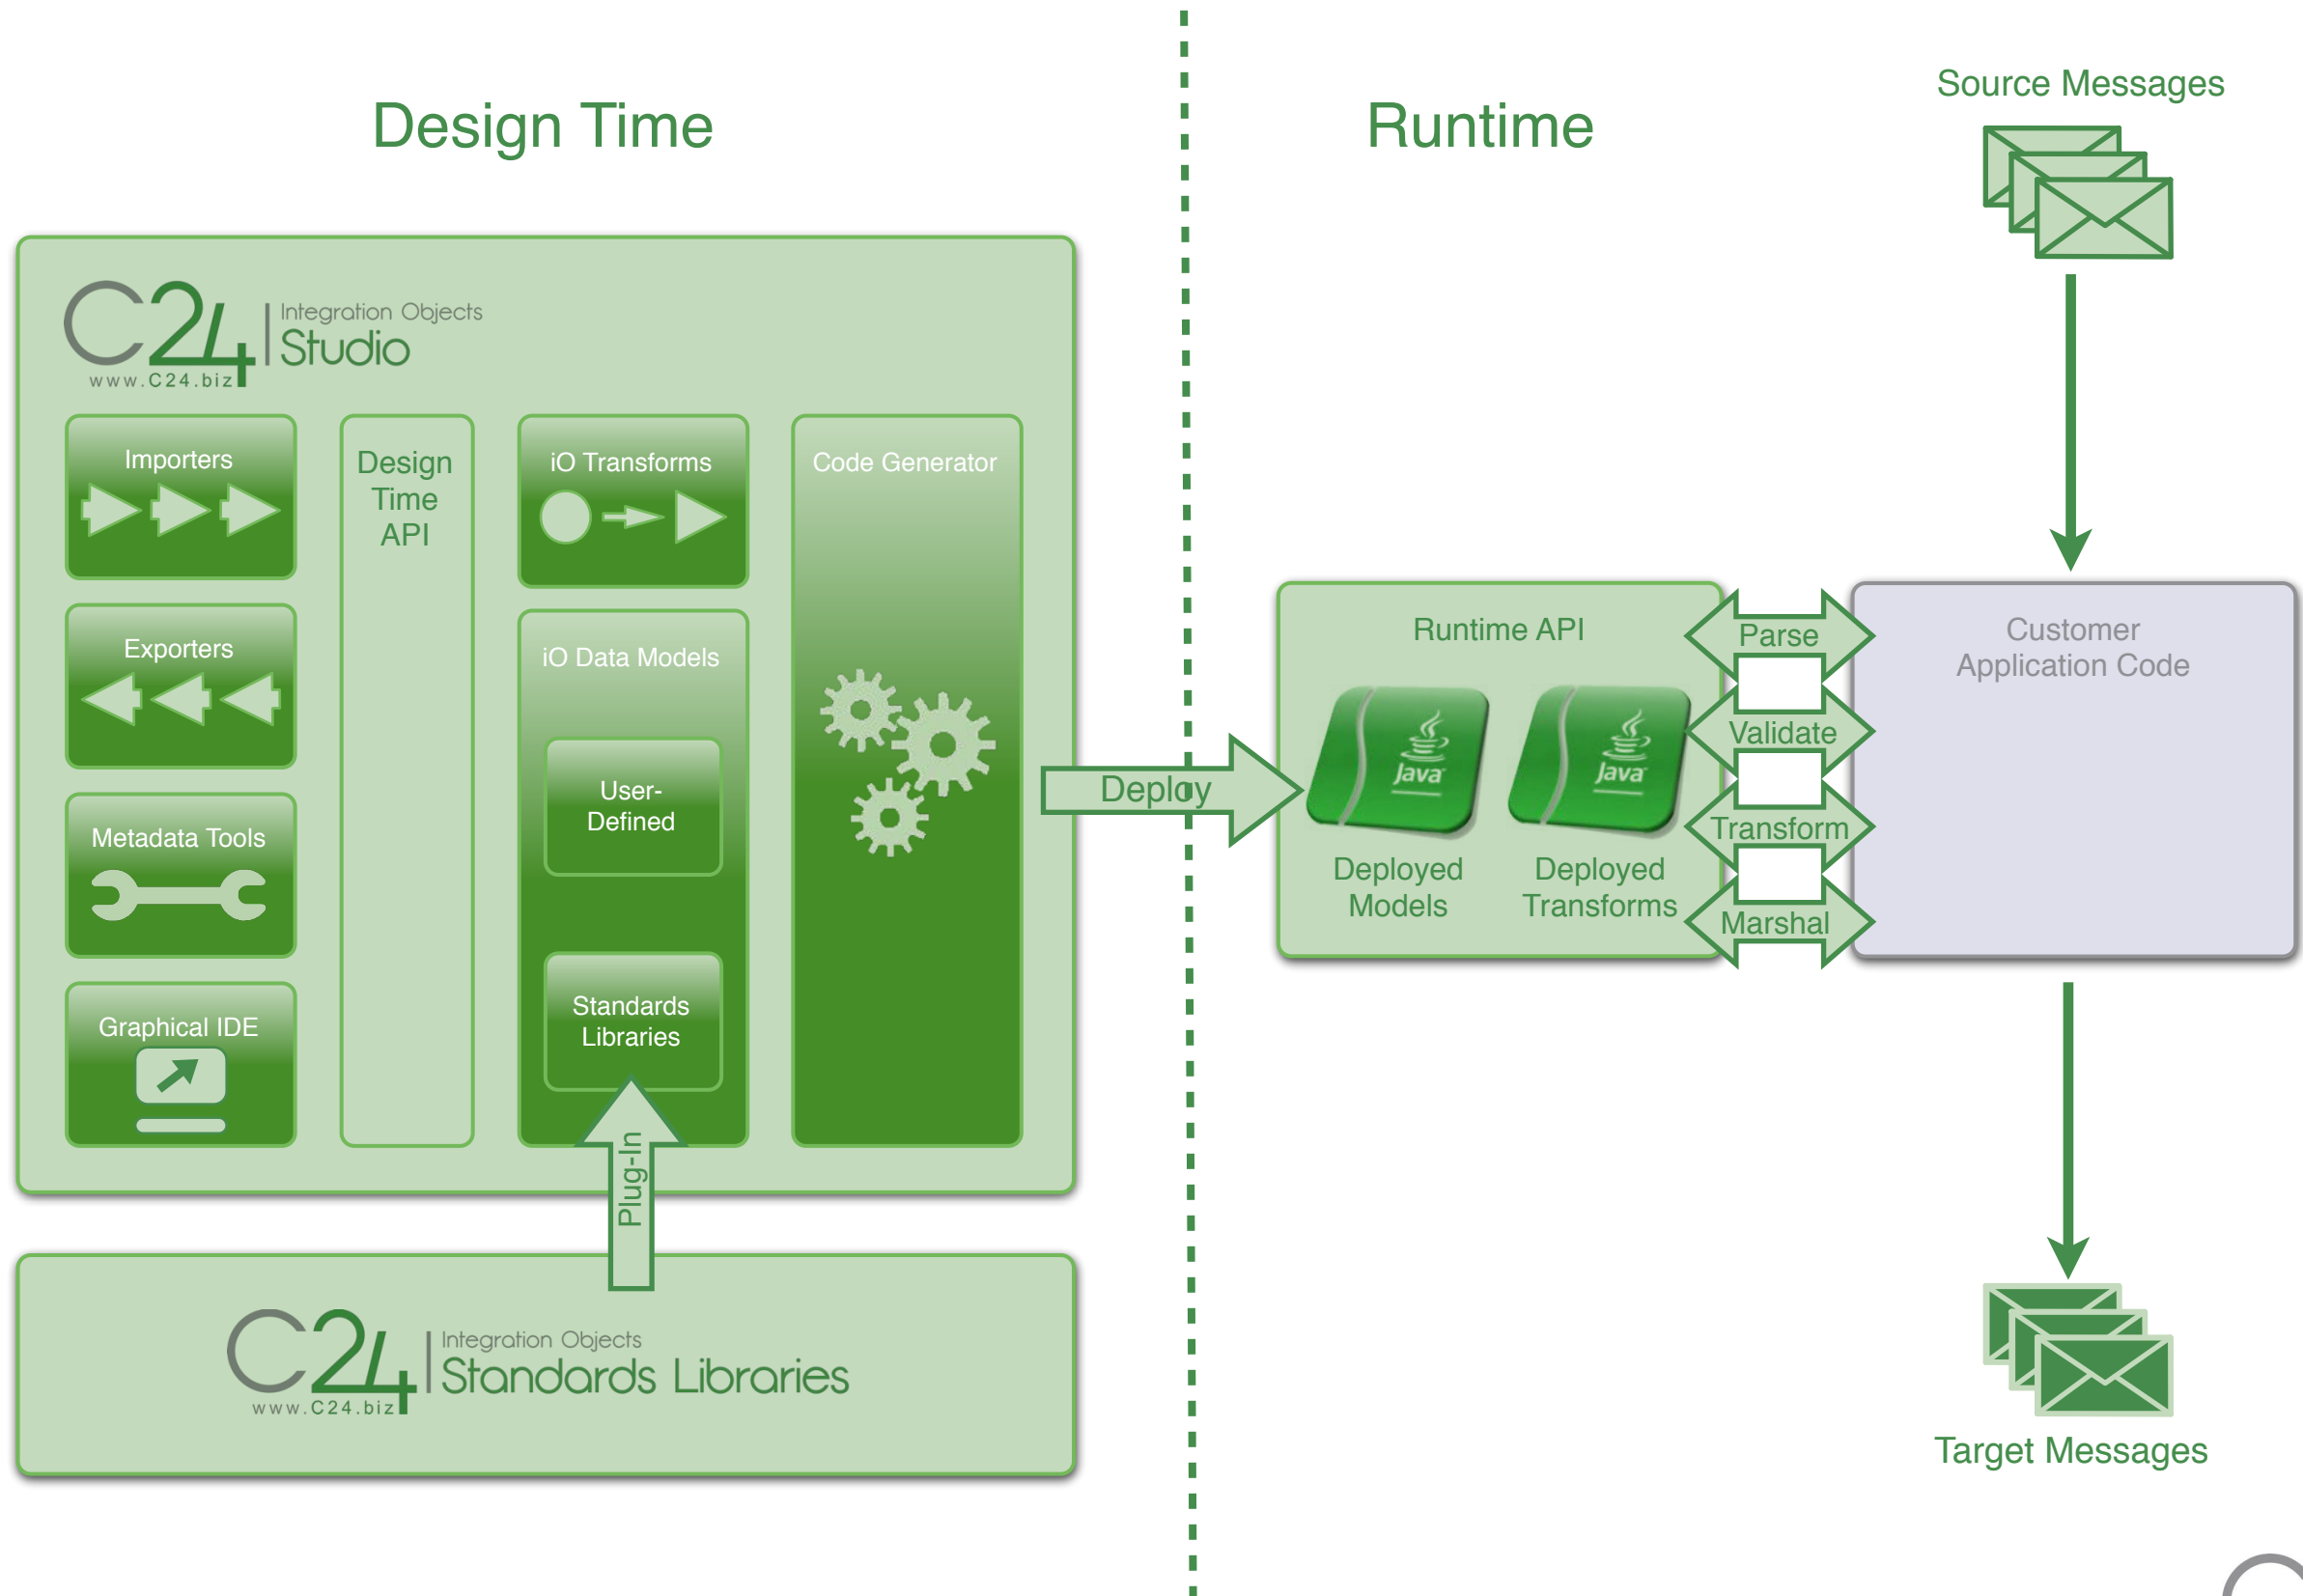
## Getting more out of (and specifically IN T0) Oracle Coherence

John Davies | CTO

Steve Miller | Product Director

Coherence SIG

Oracle London HQ | 17<sup>th</sup> July 2014





Default.asp - [Users\jaleves\My ID Projects] - ...IT FIN\SWIFT FIN November 2013\SWIFT FIN November 2013 MT564\SWIFT FIN November 2013 MT564\_3.dod - C24 Integration Objects 4.5.1 Professional Edition

MT564...

1. Project

- SWIFT FIN November 2013 MT4nn
- SWIFT FIN November 2013 MT5nn
  - Specs
  - Test Data
  - SWIFT FIN November 2013 MT5nn Comm
  - SWIFT FIN November 2013 MT5nn\_1.dod
  - SWIFT FIN November 2013 MT5nn\_2.dod
  - SWIFT FIN November 2013 MT5nn\_3.dod
  - SWIFT FIN November 2013 MT5nn\_4.dod
- SWIFT FIN November 2013 MT6nn

2. Explorer

Native

- MT564 Corporate Action Notification
  - MT564
  - MT564i
  - MT564a
  - MT564 Message
    - Block2 (local)
    - MT564 Sequence A General Information

Referenced

Built-in

3. Find

4. Properties

Properties Of **Field 98a Date - Qualifier (Read Only) - Atomic Simple Data Type**

**Presentation**

Inherent: [1a]  
Terminator: / 125  
Full:

**Presentation / Advanced**

Override Of Parse Method:  
Override Of Format Method:  
Whitespace: Preserve  
Locale: (Default)  
TLV Form Factor: None

**Validation**

Enumeration: Qualifier  
Pattern: 4N, SWIFT  
Validation Rules: 0

5. MT564 Message

Component	Type	Cardinality	Size
MT564 Message	MT564 Message		190 - "
Block2	Basic Header Block 2	1	28
Block2	User Header Block 3	1	21 - 51
Block3	MT564 Text Block	0..1	4 - 97
Block4	MT564 Sequence A General	1	140 - "
SeqA	MT564 Sequence A General	1	78 - "
Field 16R	Field 16a Type (local)	1	8 - 23
Field 28E	Field 28E Type (local)	0..1	7 - 12
Field 20a Reference	Field 20a Type37	1..*	15 - 30
C	Field 20a Reference Option C	1	10 - 25
Qualifier	Field 20a Reference -	1	6
Reference	Field 20a Reference -	1	2 - 17
Field 23G Function of the	Field 23G Type (local)	1	11 - 16
Field 22a 1	Field 22a Type94	1..*	18 - 26
Field 98a 2	Field 98a Type46	0..1	22 - 28
A	Field 98a Date Option A	1	17
Qualifier	Field 98a Date - Qualifier	1	6
Date YYYYMMDD	Field 98a Date - Date	1	9
C	Field 98a Date Option C	1	23
Field 21D 3a8 Date YYYYMMDD	Field 21D 3a8 Date YYYYMMDD	1	18 - 26
SeqA2	MT564 Sequence B	1	30 - 134
Field 16S	Field 16a Type (local)	1	8 - 23
SeqB	MT564 Sequence B	1	55 - "
SeqC	MT564 Sequence C	0..1	42 - "

Tree Graph

6. Messages

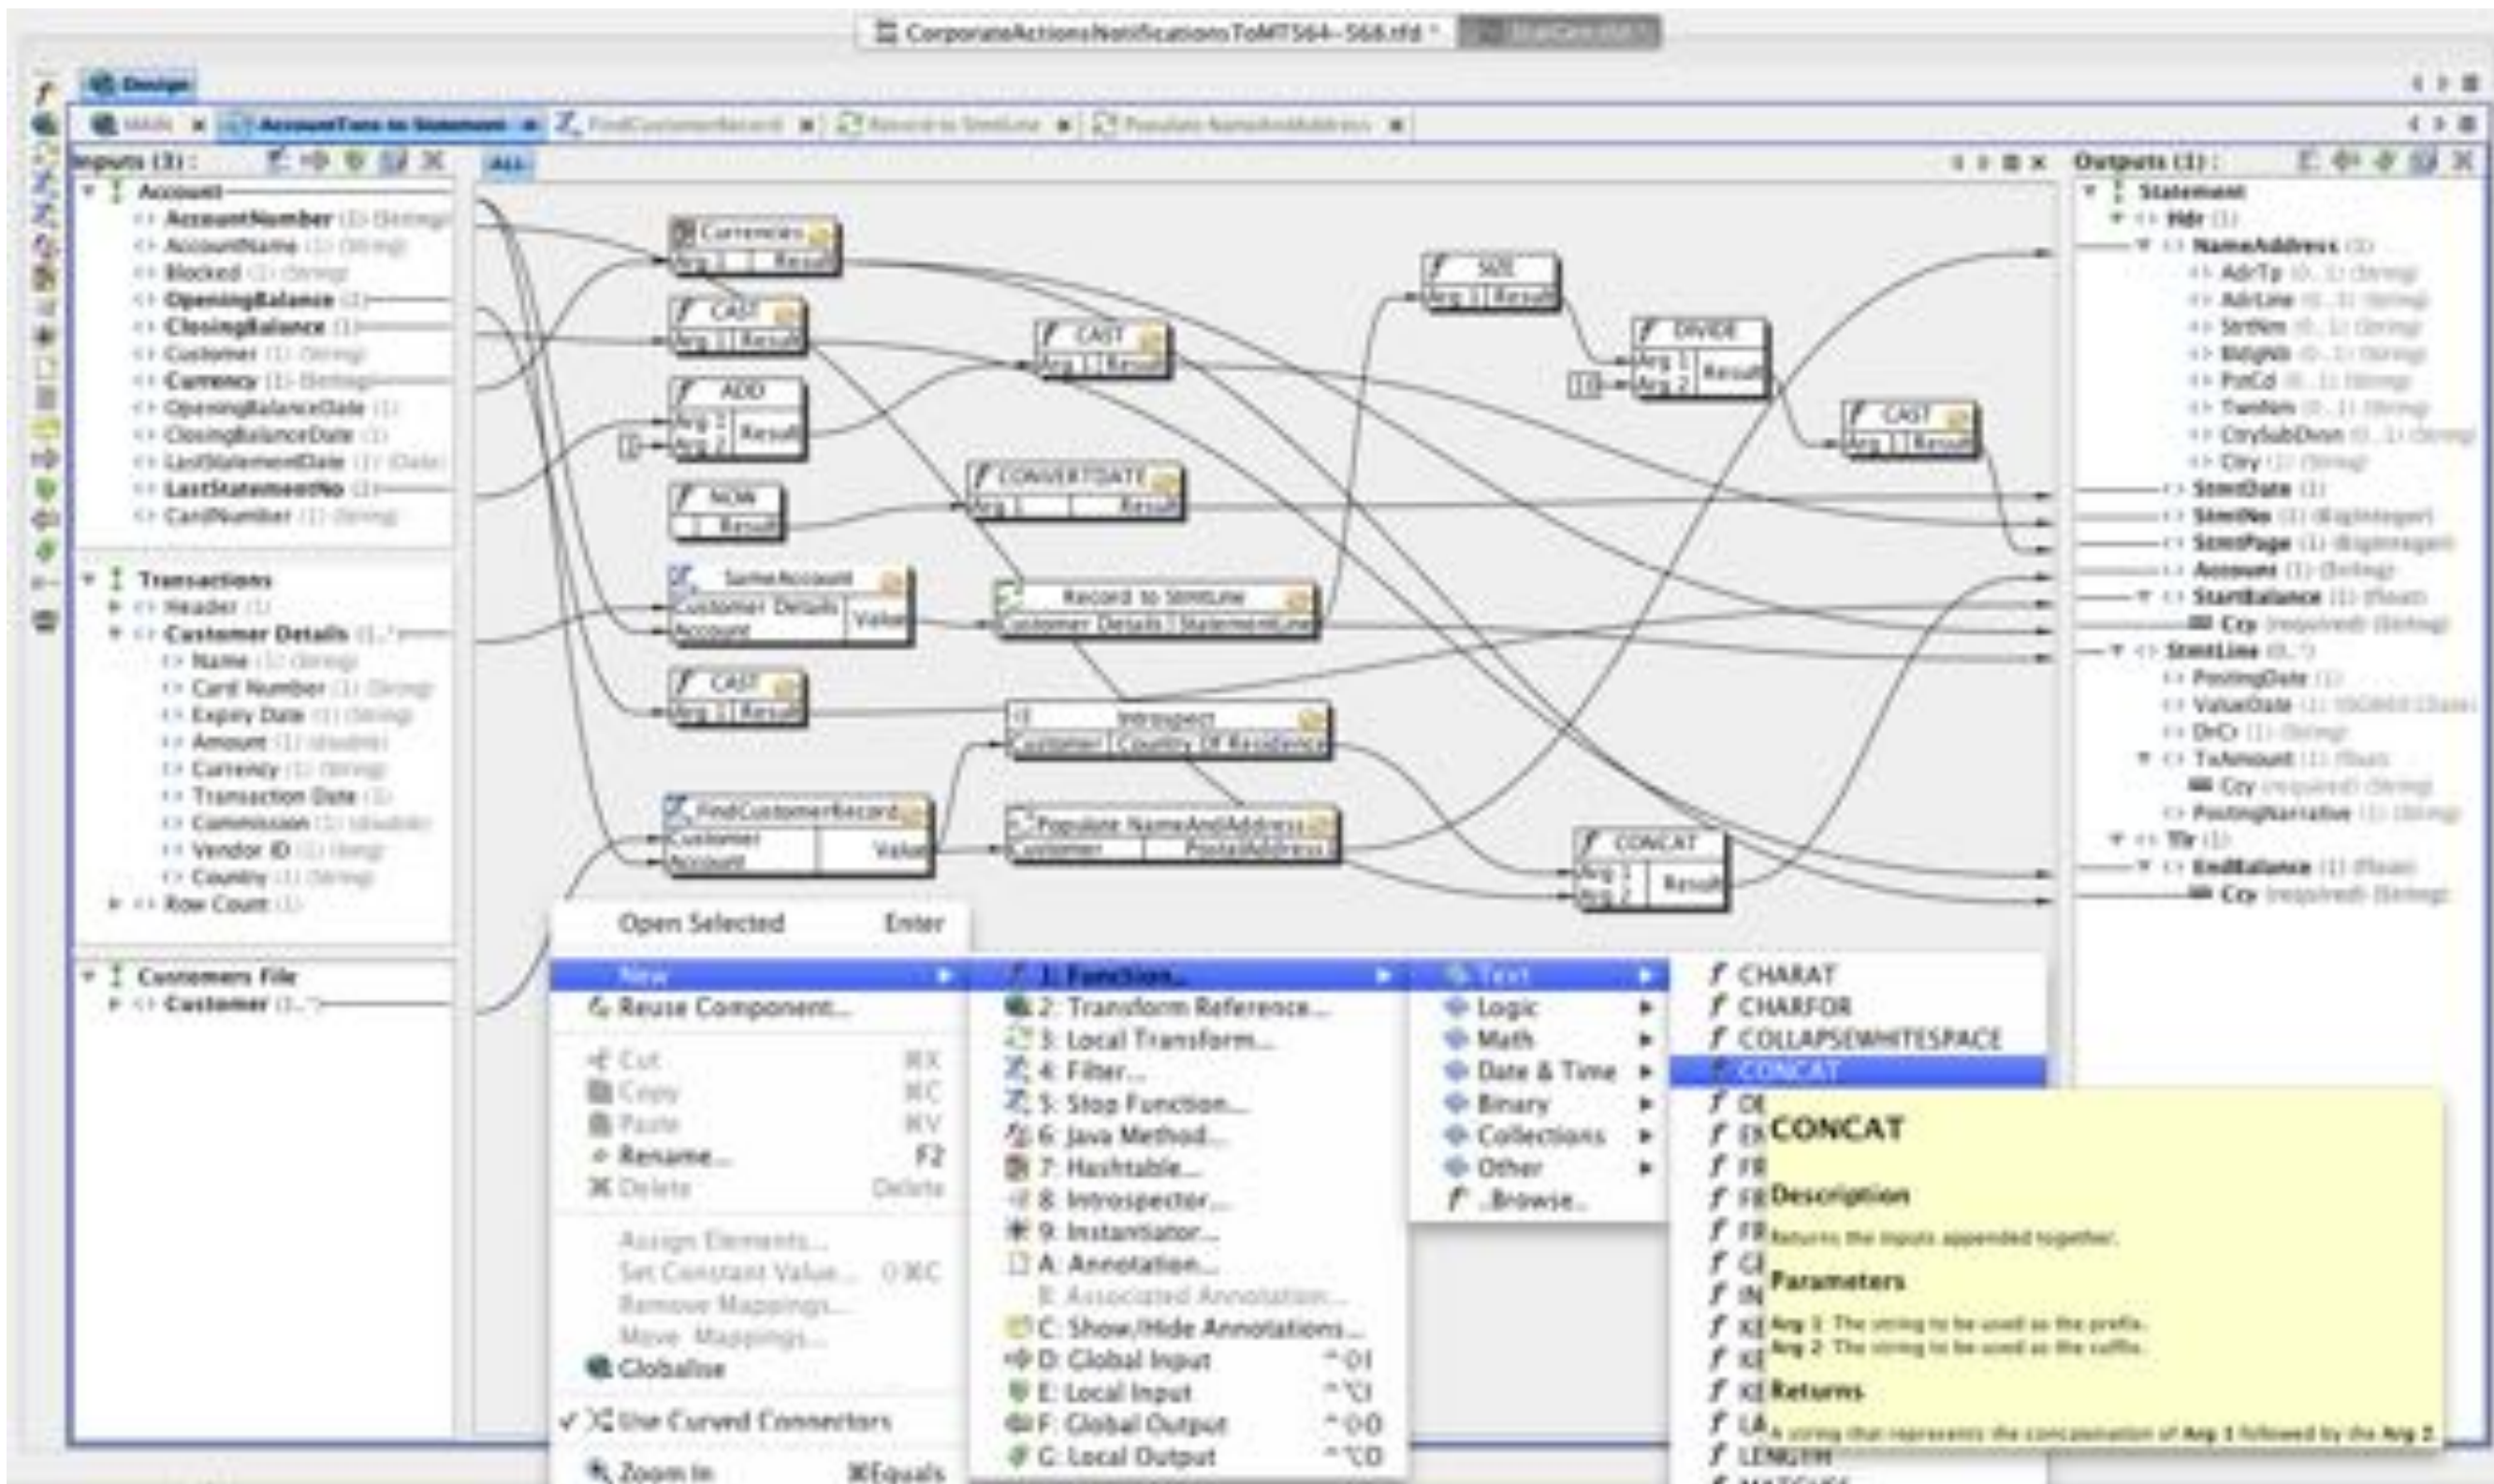
7. Properties

8. Validation

Ready

232M of 300M





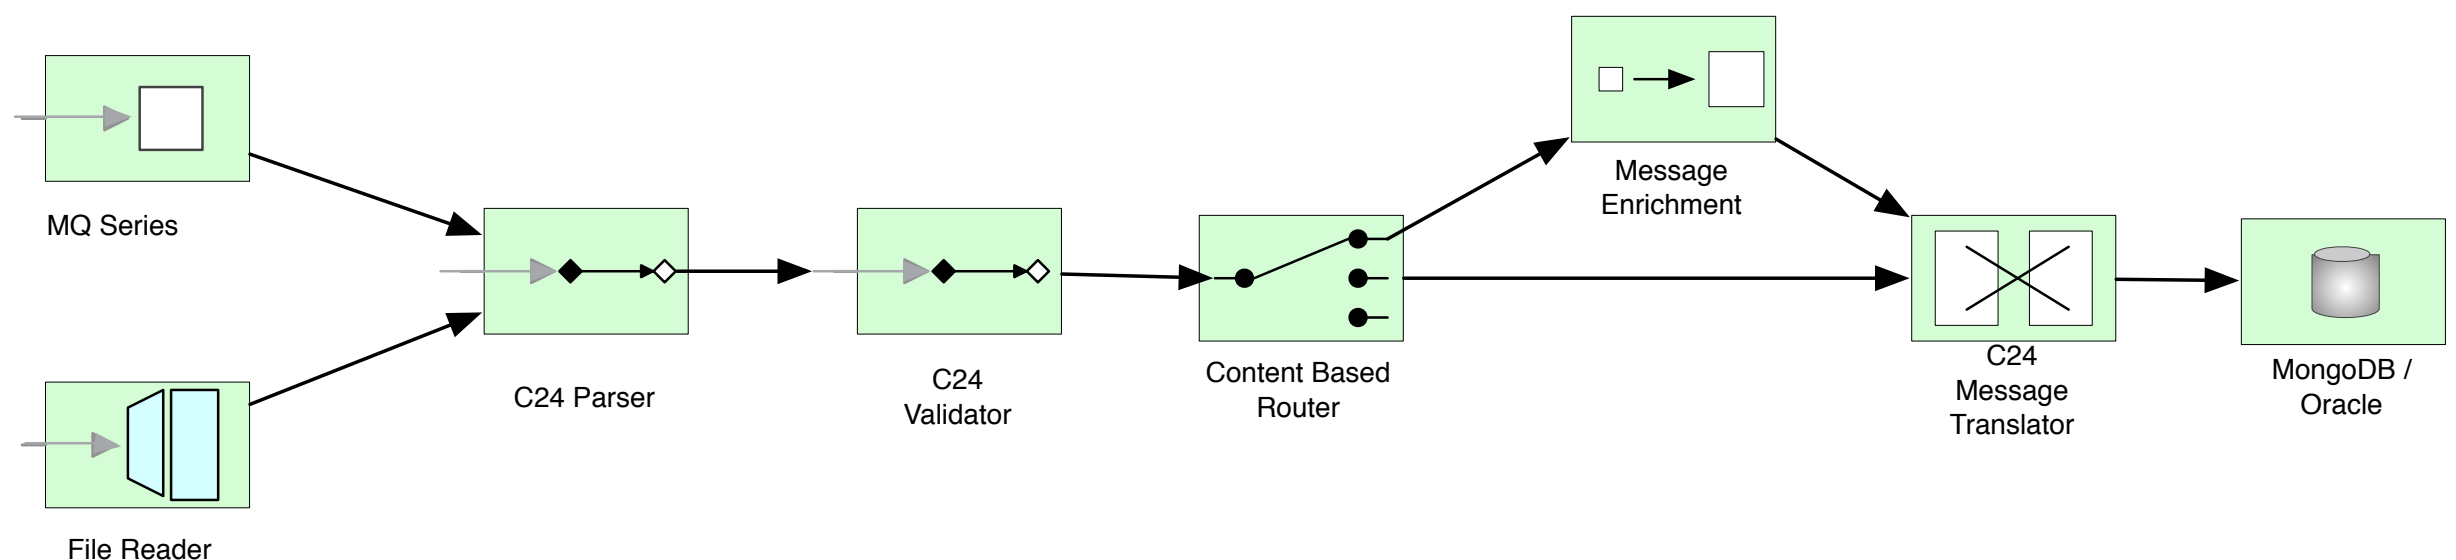


- Where ever you're dealing with events or messages, C24 can help
- The more complex or scale-critical the better the advantage
  - Proprietary formats, industry standards, legacy interfaces





- C24-iO has deep integration with Spring, Mule, Fuse & Camel etc.
- We generate the Spring config for you so you can use Spring Integration right out of the box
- New performance changes to Spring due out later this summer were driven by C24
  - Spring 4.1 will be able to pre-compile the SpEL expressions






- Deploying Fix for example will deploy the config for the parser and model...

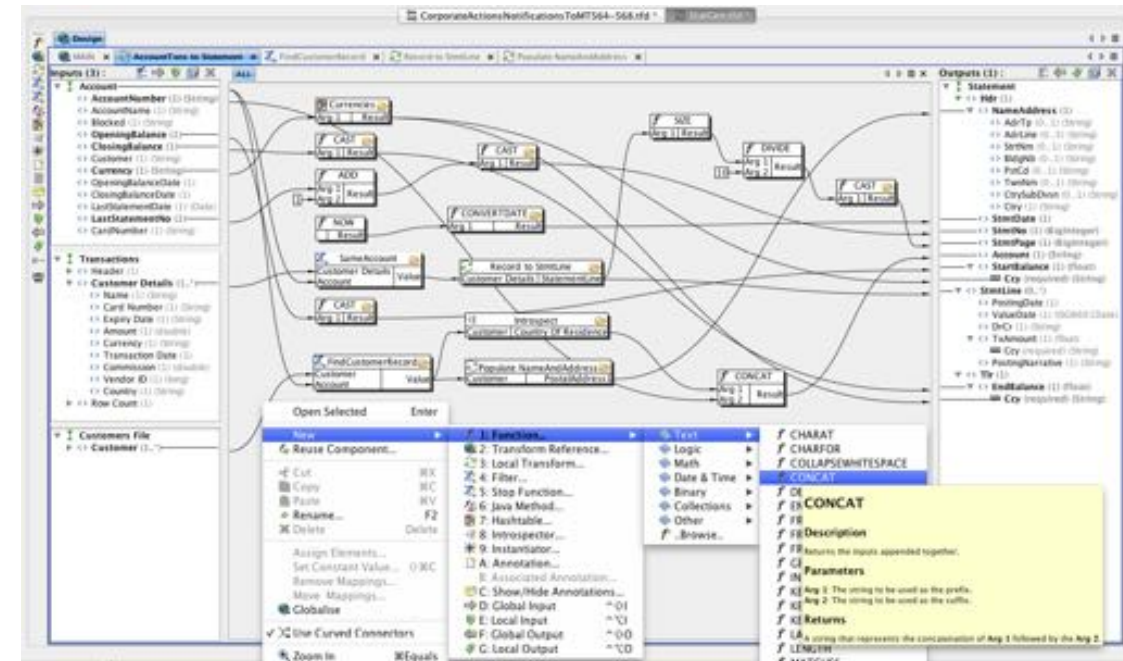
```
<int-c24:unmarshalling-transformer  
  model-ref="fixModel"  
  source-factory-ref="sourceFactory"  
  input-channel="..." output-channel="..." />
```

```
<bean id="sourceFactory"  
  class="biz.c24...source.FixSourceFactory">  
  <property name="encoding" value="UTF-8" />  
</bean>
```






- Real X-to-Y transforms are modelled in the Studio
- The generated transform slots into SI by just specifying the generated class



```
<int-c24:transformer
  transform-class="biz.c24...basic.ExampleTransform"
  input-channel="..." output-channel="..." />
```







- Change the deploy option to Java 8 and...

```
List<Atom> elements = C24
    .parse(biz.c24.periodic.PeriodicDocumentRoot.class)
    .from(new File("resources/periodic.xml"));


elements.sort((a1, a2) -> Integer.compare(a1.getAtomicNumber(), a2.getAtomicNumber()));

// List all the elements that are solid at room temperature (20C) and boil below 600 deg C
System.out.println("\nSolid at 20°C but boil lower than 800°C...");
elements.stream()
    .filter(a -> a.getMeltingPoint() != null && a.getBoilingPoint() != null )
    .filter(a -> a.getMeltingPoint().getValue() > 293 & a.getBoilingPoint().getValue() < 1073)
    .map(Atom::getNameElement)
    .forEach(System.out::println);
```

- And it works nicely in Scala too...

```
var parser = C24.parse(classOf[CustomersFile]) as C24.Format.XML
var transform = new GenerateContactListTransform
var writer = C24.write() as C24.Format.JSON

new File("/Customers.xml") -> parser -> transform -> writer -> System.out
```



- The XML and therefore generated Java API is usually technically formatted, i.e. it's not easy to extract key business data
  - tradeDate, buySideCurrency, settlementData are all hidden in the message
- For this reason most architectures use a message wrapper and extract the key fields into a header (header enrichment)
  - Even canonicalised messages present the same issue, key business fields are difficult to find
- Using C24 virtual methods...
  - No need for extra message wrappers, no extra memory used
  - Vastly simplifies user-code & maintenance
  - Can be used with ESB/SOA & messaging for filtering & routing etc.
  - Can be used with in-memory cache and database queries/QL etc.
  - Extremely powerful with Spring Integration/Mule, Coherence etc.



- We can now generate virtual getters, i.e. getters for fields that don't necessarily represent a real field in the model

```
tradeConfirmed.getTrade().getTradeHeader().getTradeDate().getValue().toDate();
```

- We can now use the much simpler...

```
Date tradeDate = tradeConfirmed.getTradeDate();
```

- Instead of this...

```
MT541SequenceE3Amounts[] seqE3 = mt541Message.getBlock4().getSeqE().getSeqE3();
for (MT541SequenceE3Amounts e3Amount : seqE3) {
    for (Field19aType31 field19 : e3Amount.getField19aAmount()) {
        if (field19.getA().getQualifier().compareTo("SETT") == 0) {
            CurrencyAmount currencyAmount = field19.getA().getSignedCurrencyAmount().getCurrencyAmount();
        }
    }
}
```

- We can use the much simpler...

```
CurrencyAmount getSettlementCurrencyAmount();
```





- Using virtual methods with lambdas we can further simplify the code

```
BigDecimal sum = transactions.stream()  
    .filter( t -> t.getBuySideCity().equals("London") )  
    .filter( t -> t.getBuySideCurrency().equals("GBP") )  
    .map( t -> t.getBuySideAmount() )  
    .sum();
```

- This now works across every version of message format and even different message formats
  - FpML
  - FIX
  - ISO 20022
  - Internal canonical format
  - CSV



- **Everything too large? Why not compress it?**
  - It's slow to compress - takes up CPU cycles
  - It's slow to decompress - takes up more CPU cycles
  - The compressed data is relatively useless until it's decompressed
  - Compressing batches is more efficient but you then have to decompress the entire batch too - More CPU cycles again
- **Compaction**
  - Smaller size but de-compaction is almost free - in some cases better
  - Works at the field level so we can use the data in its compact form
  - Compaction can use many of the features of compression
- **Take a trade value... GBP 12,500,000.00**
  - We might want to search on GBP values over 10 million
  - With compaction we can do that, compression we need to decompress first



- Typically Java Binding tools, like JAXB, JiBX and C24 create Java that looks like the data source
- While this is very convenient for the programmers it creates a lot of Java objects, this slowly consumes memory
- A typical FpML trade is around 8k in size, bind it to Java and it increases to around 25k
- 1 million FpML message in memory is going to cost anything from 8 to 25GB of RAM, add (HA) high availability and we hit 50GB
  - Expensive!
  - In-memory is still fast but 25k message over the network is very slow
  - And 25GB of data over a network or onto disk, even SSD is slow





- SDOs or Simple Data Objects are basically Java Binding into a compact binary codec - From any XML format to binary
- We analyse the data model (or XML schema) not just the instance data so can do things like...
  - Reducing the 7 days of the week to just 3 bits
  - Commonly used Strings become lookups into a static table (1 or 2 bytes)
  - Currencies for example only need 1 byte
  - Date/Time with timezone can be stored in 6 bytes
- Bit-fields are compacted resulting in excellent compaction-ratios
  - Getters calculate the offset on the fly, mask and shift the data and return it
- There is NO change to the getter API between standard binding and SDOs



```
<resetFrequency>  
  <periodMultiplier>6</periodMultiplier>  
  <period>M</period>  
</resetFrequency>
```

- JAXB, JIBX, Castor and standard C24 generate something like ...

```
public class ResetFrequency {  
    private BigInteger periodMultiplier; // Positive Integer  
    private Object period;               // Enum of D, W, M, Q, Y  
  
    public BigInteger getPeriodMultiplier() {  
        return this.periodMultiplier;  
    }  
    // constructors & other getters and setters
```

- In memory - 3 objects - at least 144 bytes
  - The parent, a positive integer and an enumeration for Period
  - 3 Java objects at 48 bytes is 144 bytes and it becomes fragmented in memory



```
<resetFrequency>
  <periodMultiplier>6</periodMultiplier>
  <period>M</period>
</resetFrequency>
```

- Using C24 SDO binary codec we generate ...

```
ByteBuffer data;    // From the root object
```

```
public BigInteger getPeriodMultiplier() {
    int byteOffset = 123;    // Actually a lot more complex
    return BigInteger.valueOf( data.get(byteOffset) & 0x1F );
}
// constructors & other getters
```

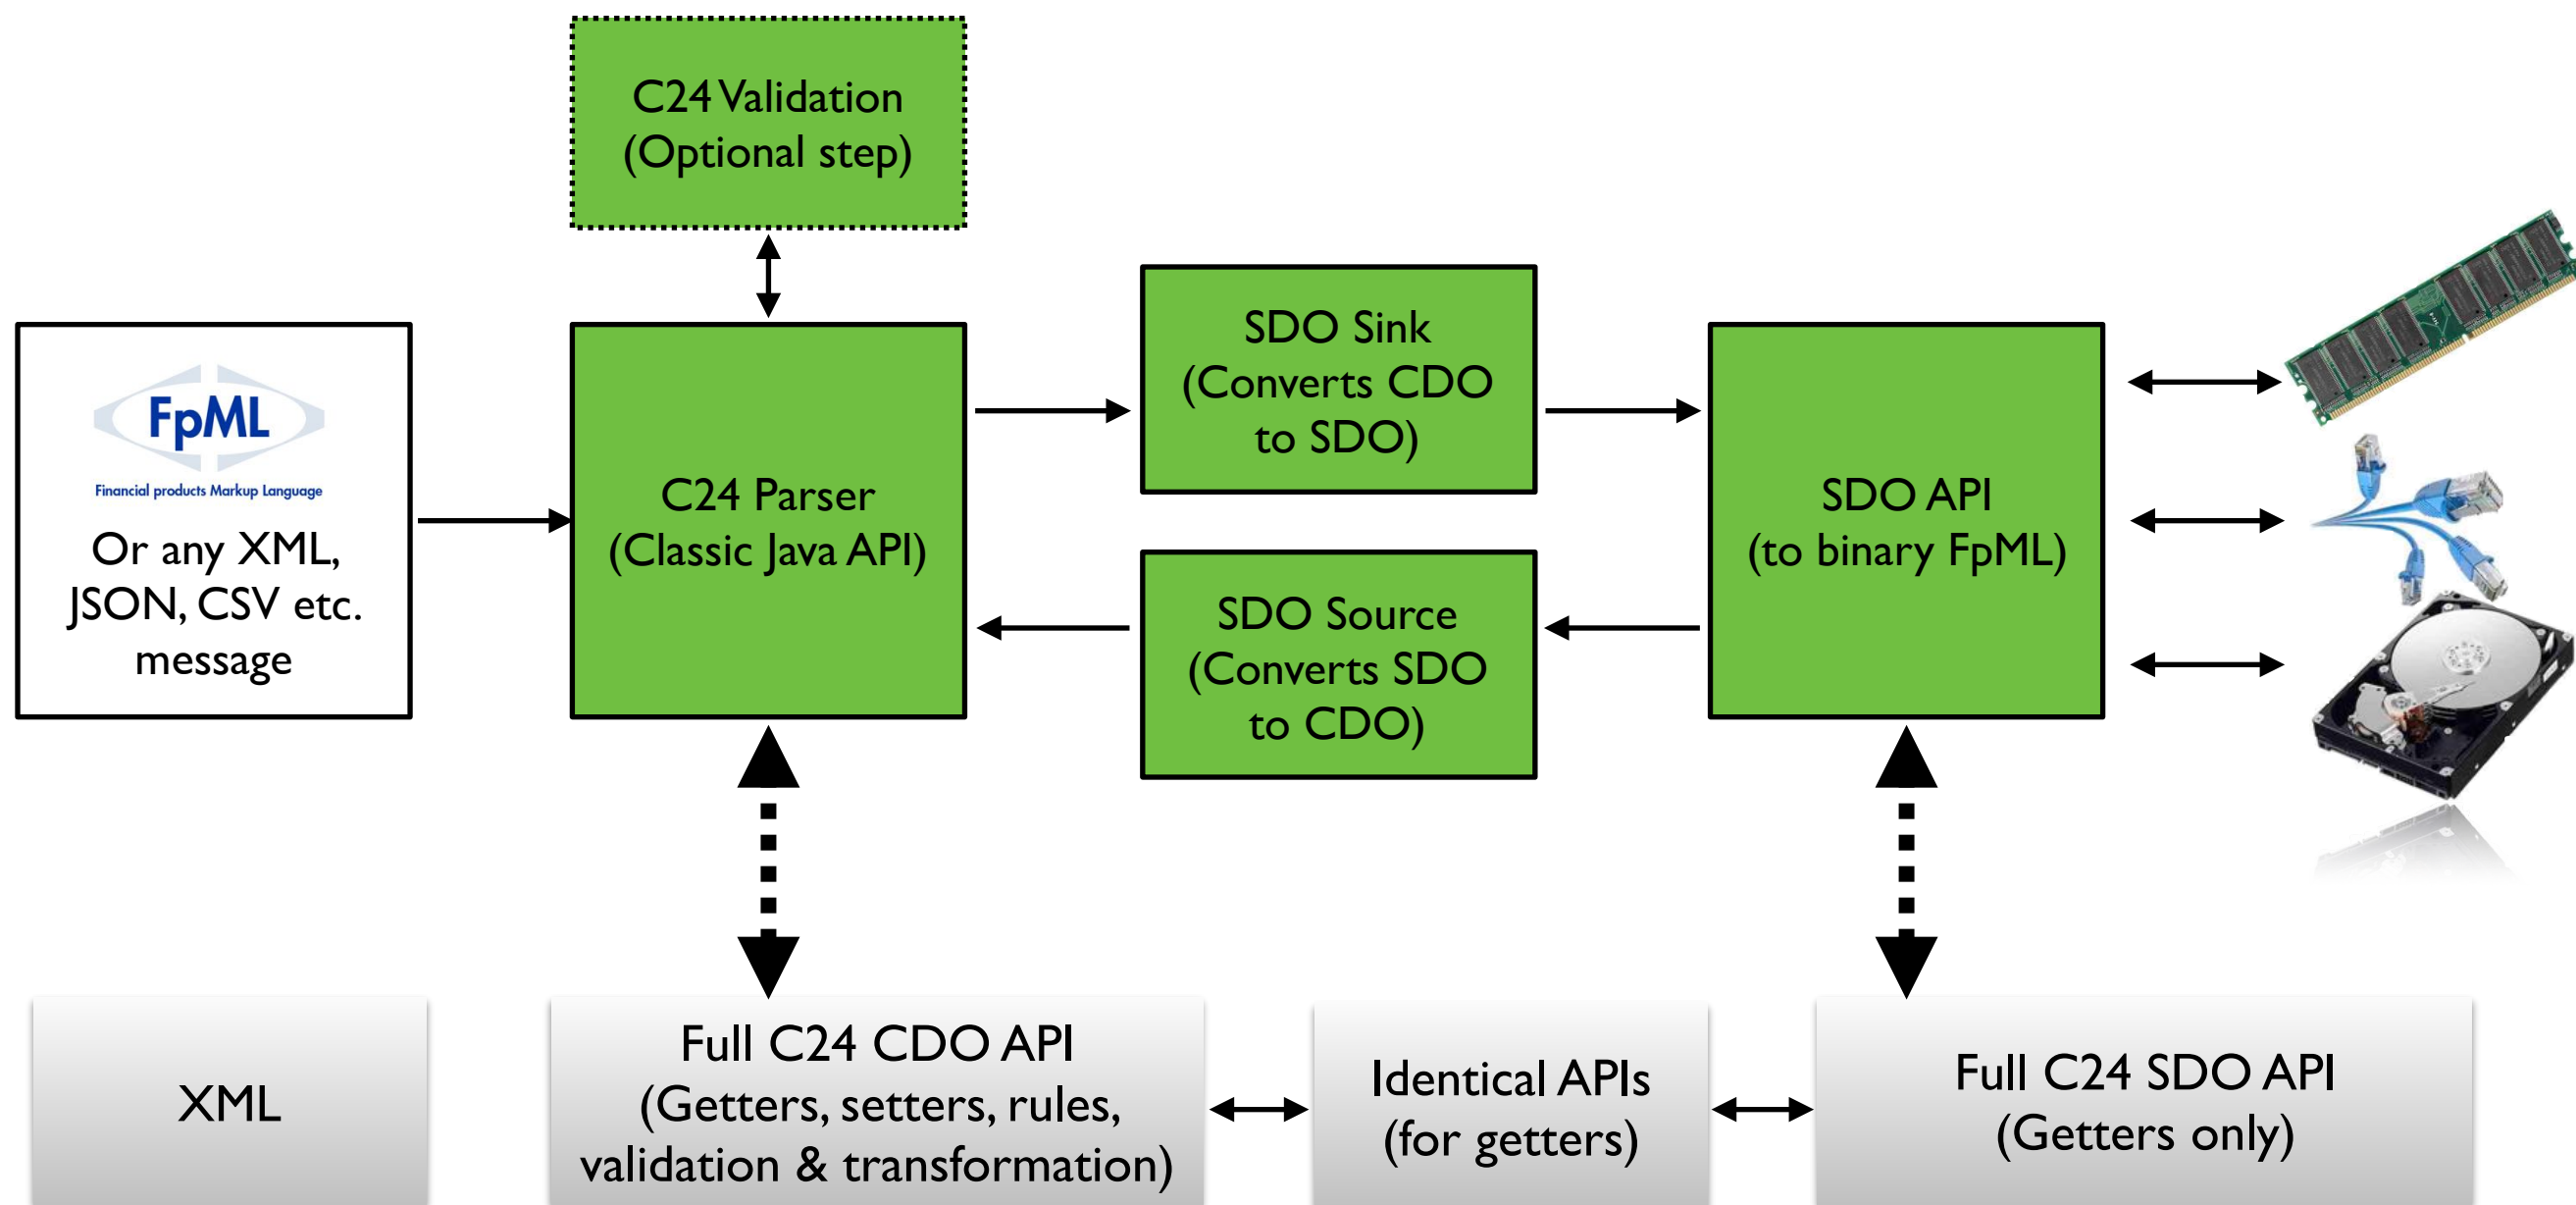
- In memory - 1 byte for all three fields
  - The root contains one ByteBuffer which is a wrapper for byte[]
  - The getters use bit-fields, Period is just 3 bits for values D, W, M, Q or Y





~5-8k	10-25k	<b>Size</b>	< 500 bytes
-------	--------	-------------	-------------

10k/sec	~1m/sec	<b>Performance</b>	~1m/sec
---------	---------	--------------------	---------





- ISDA's sample Interest Rate Derivative (vanilla swap) is 7.4k
  - We randomised a few fields and created a few million for testing
- Zipped they are average 1,547 bytes
  - 1 million on disk require 1.5GB and takes 200 seconds to read/decompress
  - Parsing at 20k/sec would add another 50 seconds and need a lot of memory
- In memory they are roughly 25k in size (in roughly 400 objects)
  - It was difficult to fit 400k into 10GB of RAM - Lots of full GCs too
- With SDOs the average size was just 442 bytes
  - It took 9 seconds to read and parse 1 million from disk (SSD)
  - It took 415ms to search through all 1 million IRSs in memory (brute force)
  - 20 million fully parsed IRSs comfortably fit in 10GB of RAM
- Total saving on memory with FpML is roughly 50 times

\* Tests were run on Java 1.7.0\_55 on a MacBook Pro (2.7 GHz Intel i7) on a single core, we continue to improve these figures



## A 5 year leap into the future with Moore's law

- In a nutshell we can compact data by typically over 10 times
- You can get at least 10 times more data into Coherence
- Data takes up a 10<sup>th</sup> of its usual size
  - On disk or in memory
- Better use of network, memory and disk
  - Massive savings in infrastructure!







Info@C24.biz  
@C24io

John.Davies@C24.biz  
@jtdavies

SDO landing page: <http://sdo.c24.biz>

<http://ref.c24.biz/whitepapers/C24-SDOs-and-Coherence.pdf>

<http://ref.c24.biz/whitepapers/C24-SDOs-Big-Data-In-Memory.pdf>